

AAYERA AHMAD

AI & Cloud Engineer · LLMs | RAG | Agentic AI | GCP | Python | FastAPI

Bengaluru, India · +91 7408478673 · aayera1029@gmail.com · linkedin.com/in/aayeraahmad531

PROFESSIONAL SUMMARY

AI & Cloud Engineer with 4 years at Harman Connected Services, working across cloud infrastructure and Generative AI development. Experienced in building LLM-powered applications, RAG pipelines, and agentic workflows using OpenAI, Gemini, LangChain, and Vertex AI. Deployed containerised AI services on GCP Cloud Run with CI/CD. Google-certified in Generative AI, Vertex AI, Agentic AI, and Responsible AI.

CORE SKILLS

AI & LLMs: OpenAI · Gemini · Vertex AI · LangChain · LangGraph · Prompt Engineering · RAG Pipelines · Agentic AI · Vector Databases · Embeddings · Model Evaluation · LLM Orchestration

Cloud & Backend: GCP (Cloud Run, Vertex AI, Cloud SQL) · Python · FastAPI · REST APIs · Docker · CI/CD (GitHub Actions)

Practices: Prompt Optimisation · Model Evaluation · Responsible AI · Agile/Scrum · JIRA

EXPERIENCE

Engineer — AI & Cloud

Harman Connected Services · Bengaluru · Mar 2022 – Present

- Developed LLM-powered applications and conversational agents using OpenAI and Gemini APIs with LangChain, improving automation and reducing manual effort across business workflows.
- Built and optimised RAG pipelines using vector databases and embedding models to ground LLM responses in verified data sources and reduce hallucinations.
- Applied prompt engineering techniques including few-shot prompting and structured outputs to improve response consistency and accuracy across production use cases.
- Engineered FastAPI microservices connecting LLM APIs and Vertex AI endpoints with internal tools and data pipelines.
- Deployed containerised AI services on GCP Cloud Run with GitHub Actions CI/CD, logging, and monitoring.
- Conducted model evaluation across OpenAI, Gemini, and Vertex AI families — assessing accuracy, latency, and cost to guide model selection decisions.
- Contributed to GCP cloud infrastructure work including Cloud Run deployments and backend API integrations supporting AI service delivery.

PROJECTS

AI Hiring Bias Detector — FastAPI · OpenAI API · Docker · GCP Cloud Run

Built a FastAPI service that uses LLM reasoning to detect gender, age, and cultural bias in job descriptions and resumes — containerised with Docker and deployed on GCP Cloud Run.

LLM Hallucination Tester — FastAPI · LangChain · OpenAI API · Docker · GCP Cloud Run

Built an automated reliability tool that generates questions on any topic, queries an LLM, and scores factual accuracy — used to benchmark model trustworthiness before deployment.

AI Policy Compliance Checker — FastAPI · OpenAI API · LangChain · Docker · GCP Cloud Run

Developed an LLM-powered compliance engine that evaluates AI-generated content against EU AI Act principles — containerised and deployed on GCP Cloud Run.

CERTIFICATIONS

- AI Agents Intensive — Agentic AI & Multi-Agent Systems — Kaggle & Google (2025)
- Introduction to Large Language Models · Generative AI with Vertex AI · Prompt Design · Responsible AI — Google Cloud (2026)
- CI/CD Pipelines on Google Cloud · Serverless Applications on Cloud Run — Google Cloud (2025)

EDUCATION

B.Tech — Information Technology

Dr. A.P.J Abdul Kalam Technical University, Lucknow · 2018 – 2022 · CGPA: 7.28